

CSC369 Tutorial: Intro to Assignment 2


Alexey Khrabrov, CSC369 Fall 2020, University of Toronto

A2 overview

- Goal: implement a message queue
 - Similar to a pipe; main difference - message boundaries
 - Main focus on synchronization
- Part 1: read/write operations
 - Classic producer-consumer plus *close* operations and *non-blocking* reads/writes
- Part 2: poll - I/O multiplexing (waiting until at least one queue is ready)
 - Modeled after the `poll()` system call
 - The hard part of the assignment
- Starter code fully specifies the API and provides an implementation skeleton
- Much “smaller” assignment compared to A1
 - Our solution adds < 300 LOC to starter code
- Difficulty stems from complexities of synchronization

Message queue

- Underlying storage - ring buffer of given capacity



- Message consists of size and contents
 - Use `size_t` to store the size
 - Contents - blob of `size` bytes
 - **Not** a null-terminated string (in general)
- Read/write operations are *atomic* and preserve message boundaries
 - Read and write one whole message (including size header) at a time
 - Write: if not enough free space for the whole message - block until enough space
 - Read: if next message is larger than user-supplied buffer - return error, leave the queue as is
 - Need to read the size without extracting it from the ring buffer

Queue handles

- API refers to a queue by a *handle*
 - Opaque identifier - similar to a file descriptor
 - `msg_queue_t msg_queue_create(size_t capacity, int flags);`
- Can open additional handles to an existing queue
 - Similar to duplicating a file descriptor
 - `msg_queue_t msg_queue_open(msg_queue_t queue, int flags);`
- Handle = reference to queue + *flags*
 - Determine what operations are allowed for this handle and their behaviour
- *Reader* handles and *writer* handles
 - `MSG_QUEUE_READER` and `MSG_QUEUE_WRITER` flags
- *Blocking* (default) and *non-blocking* handles
 - `MSG_QUEUE_NONBLOCK` flag

Close semantics

- Queue is destroyed when all handles are closed
 - Reference count reaches 0
- What if all *reader* handles are closed, but there are still *writer* handles left
 - Or the other way around
 - Similar to closing only one end of a pipe
 - NOTE: *newly created* queue with no reader (or writer) handles *yet* - different case
- If we do nothing, already blocked writers (or readers) would wait forever
- Need to notify readers when closing the last writer handle
 - “End of file” condition - like closing the write end of a pipe
- Need to notify writers when closing the last reader handle
 - “Broken pipe” condition - like closing the read end of a pipe
- “Other side” will wake up and fail the operation

Non-blocking operations

- Similar to e.g. non-blocking socket send/recv operations
 - Reads/writes on handles open with the `MSG_QUEUE_NONBLOCK` flag
 - Non-blocking write: fail with `EAGAIN` if not enough free space
 - Non-blocking read: fail with `EAGAIN` if the queue is empty
-
- NOTE: do **not** implement blocking ops as a busy loop, use condition variables

writer/producer

lock()

```
while(buffer is full){  
    wait(&empty);  
    if there is no reader > exit()  
}
```

Do the write op

signal(&full);

unlock()

reader/consumer

lock()

```
while(buffer is empty){  
    wait(&full);  
    if there is no writer > exit()  
}
```

Do the read op


signal(&empty);

unlock()

buffer_lock



ring buffer



Poll

- Purpose - wait until at least one queue in a set is ready (I/O multiplexing)
- Data structures
 - Queue: keep track of the threads currently blocked in poll calls for this queue
 - Polling thread: keep track of the current state of the queues it's waiting for to become ready
- Implementation outline
 - 1. Subscribe to events (so that other threads know when to notify this thread)
 - 2. Check if any events already triggered, otherwise block
 - 3. Check what events were triggered, return info to caller
- Note: reports what queues are ready *right now* - shortly before call returns
 - State can change by the time thread calls read or write
- Hints
 - Condition variable to wait on can be created dynamically for this poll call
 - Can use a linked list for the “wait queue” - threads waiting for a queue to become ready

Starter code overview

- msg_queue.h - API definitions and documentation/specification
- msg_queue.c - API implementation
 - NOTE: this is the **only** file you will modify
- ring_buffer{.h, .c} - underlying ring buffer implementation
- sync{.h, .c} - synchronization primitives
 - Wrappers around pthread functions - allow us to intercept them for testing purposes
- list.h - doubly-linked list implementation (useful for poll wait queue)
- errors{.h, .c} - error logging helpers
- mutex_validator{.h, .c} - synchronization debugging/testing tool
- prodcon.c, multiprod.c - example programs/tests
- Makefile - the only thing to change is compile flags (during development)

Linked list

- Can be (and should be) used to implement the wait queue for poll
- “Embedded” entries - different from linked lists you are used to
 - ```
struct s {
 ...data fields...
 list_entry entry; // stores next and prev pointers
};
```
- To get pointer to containing struct from pointer to entry:
  - ```
struct_ptr = container_of(entry_ptr, s, entry);
```
- More efficient than having separate struct with next, prev, data fields
 - Not wasting memory on storing data pointers
 - Avoiding lots of small dynamic memory allocations

Synchronization debugging tools

- **Mutex validator**
 - Checks that operations that must be done inside a critical section are actually done by one thread at a time
 - Adds delays to extend duration of critical sections to make violations more likely
 - Included in ring buffer and linked list
 - You can use it in your data structures for implementing poll
- **Early wakeups from condition variable waits**
 - Can be useful in debugging cases when waiting thread “misses” a signal and waits forever
 - See details in `sync.c`
- **External tools**
 - Helgrind (part of Valgrind) - able to detect various types of synchronization error

What to implement in msg_queue.c

- Part 1

- `struct mq_backend, mq_init(), mq_destroy()`
 - Synchronization variables and their initialization
- `msg_queue_write()`
- `msg_queue_read()`
- `msg_queue_close()`
 - Synchronization and notifying the other end if closing the last reader or writer handle

- Part 2

- `struct mq_backend, mq_init(), mq_destroy()`
 - Data structures to keep track of threads blocked in poll for this queue
- `msg_queue_poll()`
- `msg_queue_close()`
 - Notifying polling threads if closing the last reader or writer handle

Important notes

- Read the handout and doc comments in starter code carefully a few times
- Common mistakes
 - Not preserving atomicity of reads and writes
 - Use `ring_buffer_peek()` to read message size
 - Incorrect use of condition variables; should always do this:

```
while (!predicate) {  
    ...check if still need to wait...  
    cond_wait(&cond, &mutex);  
}
```
 - Thread waiting forever in poll after missing a condition variable signal
- Most of error checking is about API misuse - programmers' mistakes
- Use macros from `errors.h` in the starter code for error logging